

# CSS 262: Linux Administration

Process Management & Systemd

Lecture 3: Mastering Processes & Services





# Today's Agenda

## Part 1: Process Management

- What are processes?
- Process lifecycle & states
- Process monitoring commands
- Managing processes (signals, jobs)
- Process priorities & nice values
- `/proc` filesystem

## Part 2: Systemd

- Introduction to systemd
- Unit files & service management
- `systemctl` commands
- Creating custom services
- Boot process & targets
- Logs with `journalctl`

🎯 **Learning Objective:** Master process control and modern service management for production systems.



# Quick Recap: Week 2

## What We Covered

- User and group management
- File ownership and permissions (rwx)
- chmod, chown, chgrp commands
- Special permissions (SUID, SGID, sticky bit)
- Access Control Lists (ACLs)

## Key Takeaway

Proper permission management is the foundation of Linux security - principle of least privilege.

 **Assumption:** You can now create users, manage groups, and set appropriate file permissions.

# Part 1: Process Management



Understanding how Linux manages running programs

# What is a Process?

## Definition

A **process** is an instance of a running program. Every command you run, every application you open, creates at least one process.

## Key Characteristics

- **PID:** Unique Process ID
- **PPID:** Parent Process ID
- **Owner:** User running the process
- **Memory:** Allocated RAM
- **State:** Running, sleeping, stopped, etc.

## Process vs Program

- **Program:** Static file on disk
- **Process:** Program in execution
- One program can have multiple processes
- Example: Multiple Firefox windows



**Key Concept:** Everything in Linux runs as a process - from your shell to system services.

# Process Hierarchy

## Linux Process Tree

All processes descend from **init** (PID 1) - the first process started by the kernel.

```
1  systemd (PID 1)
2  └── sshd (PID 856)
3    └── sshd (PID 2341) - user connection
4      └── bash (PID 2342)
5        └── vim (PID 2450)
6  └── apache2 (PID 1024)
7    ├── apache2 (PID 1025) - worker
8    └── apache2 (PID 1026) - worker
9  └── cron (PID 789)
```

## Orphan & Zombie Processes

- **Orphan:** Parent dies → adopted by init
- **Zombie:** Child exits but parent hasn't read status → `<defunct>`

# Viewing Processes: ps

## The Classic Process Snapshot Command

```
1 # Show your processes
2 ps
3
4 # Show all processes (BSD style)
5 ps aux
6
7 # Show process tree
8 ps auxf
9
10 # Show processes for specific user
11 ps -u john
```

## Understanding ps aux Output

```
1 USER PID %CPU %MEM    VSZ   RSS TTY STAT START TIME COMMAND
2 john 1234 2.5 1.3 123456 8192 pts/0 S+ 10:00 0:02 python script.py
```

**Key Columns:** USER (Owner), PID (Process ID), %CPU/%MEM (Usage), STAT (State), COMMAND (Program)

# Process States

## STAT Column in ps aux

Code	State	Description
R	Running	Executing on CPU
S	Sleeping	Waiting for event
D	Disk Sleep	Waiting for I/O
T	Stopped	Suspended (Ctrl+Z)
Z	Zombie	Terminated but not reaped
< / N	Priority	Nice < 0 / Nice > 0
+	Foreground	In foreground group

Examples: S+ , R< , Ss

# Dynamic Process Monitoring: top

## Interactive Process Viewer

```
1  top          # Launch top  
2  top -u john    # Show only john's processes  
3  top -p 1234,5678 # Monitor specific PIDs
```

## Top Display (First Few Lines)

```
1  top - 10:23:45 up 5 days, load average: 0.52, 0.48, 0.45  
2  Tasks: 245 total, 1 running, 244 sleeping  
3  %Cpu(s): 12.5 us, 3.2 sy, 83.8 id, 0.3 wa  
4  MiB Mem: 15924 total, 2341 free, 8456 used  
5  MiB Swap: 2048 total, 2048 free, 0 used
```

## Interactive Commands Inside top

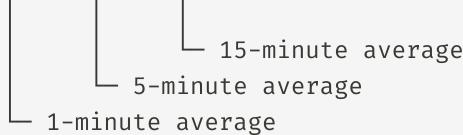
- `k` - Kill process | `r` - Renice | `u` - Filter by user
- `M` - Sort by memory | `P` - Sort by CPU | `T` - Sort by time
- `1` - Show individual CPUs | `q` - Quit

 **Pro Tip:** Use 'htop' for better UX

# Understanding Load Average

## The Three Numbers

```
1 load average: 0.52, 0.48, 0.45
2
3
4
5
```



## What Does It Mean?

Load average = number of processes waiting for CPU time.

### Rule of Thumb (for a 4-core system):

- < 4.0 - System has capacity
- = 4.0 - System fully utilized
- > 4.0 - System overloaded

**Example:** N-core system at Load N.0 = 100% utilized

# Modern Alternative: htop

## Enhanced Process Viewer

```
1 sudo apt install htop      # Install  
2 htop                      # Launch
```

### Advantages over top

- Color-coded output
- Mouse support
- Visual CPU/memory bars
- Tree view built-in
- Easier to kill processes
- Scroll horizontally/vertically

### Keyboard Shortcuts

- F9 - Kill process
- F7/F8 - Adjust nice value
- F5 - Tree view
- F6 - Sort by column
- F4 - Filter processes
- / - Search
- Space - Tag multiple

⌚ **Recommendation:** Use htop for interactive monitoring, ps/top for scripting.

# Process Signals

Signals are Messages Sent to Processes

```
1 kill -TERM 1234 # Send SIGTERM
2 kill -9 1234     # Force kill (SIGKILL)
3 killall firefox # Kill by name
```

**Common Signals:** SIGHUP (1), SIGINT (2), SIGKILL (9), SIGTERM (15), SIGSTOP (19)

# Killing Processes Properly

## The Right Way to Stop Processes

```
1 # Step 1: Try graceful shutdown
2 kill 1234
3 kill -TERM 1234      # Same as above (TERM is default)
4
5 # Wait 5-10 seconds ...
6
7 # Step 2: If still alive, force kill
8 kill -9 1234
9 kill -KILL 1234      # Same as above
```

### Good Practice ✓

1. Try SIGTERM first
2. Wait for graceful shutdown
3. Use SIGKILL as last resort

### Avoid ✗

- Starting with `kill -9`
- Killing init (PID 1)
- Killing critical system processes

# Foreground vs Background Jobs

## Job Control in the Shell

```
1 # Run in background
2 ./script.sh &
3
4 # Suspend running process (Ctrl+Z), then resume
5 bg          # Resume in background
6 fg          # Resume in foreground
7
8 # Job control
9 jobs        # List jobs
10 fg %2      # Bring job 2 to foreground
11 bg %1      # Send job 1 to background
```

# Process Management Examples

## Real-World Scenarios

```
1 # Find and kill all Python processes
2 pkill -f python
3
4 # Find process using specific port
5 lsof -i :8080
6
7 # Kill process using port 8080
8 kill $(lsof -t -i :8080)
9
10 # See what files a process has open
11 lsof -p 1234
```

# Process Priority: nice & renice

## Controlling Process CPU Priority

**Nice values:** -20 (highest priority) to 19 (lowest priority) • **Default:** 0

```
1 # Start with low priority (nice)
2 nice -n 10 ./cpu-intensive-task.sh
3
4 # Start with high priority (requires root)
5 nice -n -10 ./important-task.sh
6
7 # Change priority of running process
8 renice -n 5 -p 1234          # Set nice to 5
9 renice -n 10 -u john         # All john's processes
10
11 # View nice values
12 ps -eo pid,ni,cmd           # Show PID, nice, command
```

 **Use Case:** Run backups with `nice 19`

# The /proc Filesystem

Virtual filesystem exposing kernel and process data.

## Examples:

- `/proc/1234/` - Process-specific info
- `/proc/cpuinfo` - System-wide info
- `/proc/meminfo` - Memory details

# Useful Process Commands

Command	Purpose
<code>ps</code> / <code>top</code> / <code>htop</code>	View/monitor processes
<code>pgrep</code> / <code>kill</code>	Find/terminate PID
<code>nice</code> / <code>renice</code>	Adjust priority

# Part 2: Systemd



Modern Linux service management and init system

# What is Systemd?

## The Modern Init System

**Systemd** is the init system (PID 1) used by most modern Linux distributions. It manages system boot, services, and system state.

### Replaced

- **SysVinit** (traditional init scripts)
- **Upstart** (Ubuntu's previous init)

### Key Features

- Parallel service startup
- On-demand activation
- Service dependencies
- Resource control (cgroups)

### Systemd Components

- `systemctl` - Service manager
- `journalctl` - Log viewer
- `systemd-analyze` - Boot analysis
- `hostnamectl` - Hostname management
- `timedatectl` - Time/date settings

# Systemd Units

Everything in Systemd is a Unit

Unit Type	Extension	Purpose
Service	.service	System services
Socket	.socket	IPC sockets
Target	.target	Group of units
Timer	.timer	Scheduled tasks

## Unit File Locations:

1. /etc/systemd/system/ - Admin units
2. /lib/systemd/system/ - Package units

# Basic Service Management: systemctl

## Essential Commands

```
1  # Service Status
2  systemctl status nginx          # Detailed status
3  systemctl is-active nginx       # Check if running
4  systemctl is-enabled nginx      # Check if starts at boot
5
6  # Start/Stop/Restart
7  systemctl start nginx          # Start service
8  systemctl stop nginx           # Stop service
9  systemctl restart nginx        # Stop, then start
10 systemctl reload nginx         # Reload config (if supported)
11 systemctl reload-or-restart nginx # Reload if possible, else restart
12
13 # Enable/Disable (boot time)
14 systemctl enable nginx         # Start at boot
15 systemctl disable nginx        # Don't start at boot
16 systemctl enable --now nginx   # Enable and start immediately
17
18 # View Configuration
19 systemctl cat nginx            # Show unit file
20 systemctl show nginx           # Show all properties
21 systemctl list-dependencies nginx # Show dependencies
```

# List Services

## Finding and Filtering Units

```
1 # List services
2 systemctl list-units --type=service
3 systemctl list-units --type=service --all    # Including inactive
4
5 # List failed/enabled
6 systemctl --failed
7 systemctl list-unit-files --state=enabled
8
9 # Search
10 systemctl list-units | grep ssh
```

**Pro Tip:** Use `systemctl --failed` to find problems

# Understanding Service Status

## Reading systemctl status Output

```
1  systemctl status nginx

1 ● nginx.service - A high performance web server
2       Loaded: loaded (/lib/systemd/system/nginx.service; enabled)
3         Active: active (running) since Mon 2026-02-10 10:30:15 UTC
4           PID: 1234 (nginx)
5             Tasks: 5 (limit: 4620)
6             Memory: 12.5M
```

## Key Information

- **Loaded:** Unit file location and boot status
- **Active:** Current state (active, inactive, failed)
- **Main PID:** Primary process ID
- **Memory/CPU:** Resource usage

# Service Unit File Structure

## Anatomy of a .service File

```
1 [Unit]
2 Description=My Custom Application
3 After=network.target          # Start after network
4 Requires=postgresql.service   # Dependency required
5
6 [Service]
7 Type=simple                  # Service type
8 User=appuser                 # Run as user
9 WorkingDirectory=/opt/myapp
10 Environment="NODE_ENV=production"
11 ExecStart=/usr/bin/node server.js
12 ExecReload=/bin/kill -HUP $MAINPID # Reload command
13 Restart=on-failure           # Auto-restart on crash
14
15 [Install]
16 WantedBy=multi-user.target   # Enable in this target
```

# Service Types

Type= Directive Options

Type	Description	Use Case
<b>simple</b>	Main process is ExecStart	Foreground services
<b>forking</b>	Process forks, parent exits	Daemons (nginx, apache)
<b>oneshot</b>	Process exits after starting	Setup scripts

Examples:

- `Type=simple` - Process stays in foreground
- `Type=forking` - Traditional daemon with PIDFile
- `Type=oneshot` - Run once and exit

# Creating a Custom Service

## Example: Python Web Application

### 1. Create service file: /etc/systemd/system/myapp.service

```
1 [Unit]
2 Description=My Python Web Application
3 After=network.target
4
5 [Service]
6 Type=simple
7 User=www-data
8 WorkingDirectory=/opt/myapp
9 ExecStart=/opt/myapp/venv/bin/python app.py
10 Restart=always
11
12 [Install]
13 WantedBy=multi-user.target
```

### 2. Activate:

```
1 sudo systemctl daemon-reload
2 sudo systemctl enable --now myapp.service
```

# Systemd Targets

## Targets = Runlevels in Systemd

Targets define system states (similar to SysV runlevels).

Target	Description
poweroff.target	Shutdown
rescue.target	Single-user mode
multi-user.target	Multi-user, no GUI
graphical.target	Multi-user with GUI
reboot.target	Reboot

### Commands:

```
1  systemctl get-default  
2  sudo systemctl set-default multi-user.target
```

# Systemd Timers

Modern alternative to cron for scheduled tasks.

## Files needed:

- `backup.service` - Service definition with `ExecStart`
- `backup.timer` - Timer with `OnCalendar=*-*-* 02:00:00`

**Activate:** `systemctl enable --now backup.timer`

# Viewing Logs: journalctl

Systemd's unified logging system.

## View logs:

- `journalctl` - All logs
- `journalctl -f` - Follow

## Filter:

- `journalctl -u SERVICE` - By service
- `journalctl --since "1 hour ago"` - By time

# Advanced journalctl

## More Powerful Log Analysis

```
1 # Kernel messages
2 journalctl -k
3
4 # Multiple services
5 journalctl -u nginx -u apache2 -f
6
7 # Limit output
8 journalctl -n 50          # Last 50 lines
9
10 # Disk management
11 journalctl --disk-usage
```

# Boot Analysis: systemd-analyze

## Analyze System Boot Performance

```
1 # Overall boot time
2 systemd-analyze
3 # Output: Startup finished in 2.5s (kernel) + 8.3s (userspace) = 10.8s
4
5 # Service breakdown & critical path
6 systemd-analyze blame
7 systemd-analyze critical-chain
8
9 # Generate boot chart & verify units
10 systemd-analyze plot > boot-chart.svg
11 systemd-analyze verify /etc/systemd/system/myapp.service
```

💡 Use `systemd-analyze blame` to find slow services

# Resource Control with Systemd

## Limits Service Resources

```
1 [Service]
2 # CPU limits
3 CPUQuota=50%          # Max 50% of one CPU
4
5 # Memory limits
6 MemoryMax=512M        # Hard limit
7
8 # Task limits
9 TasksMax=100           # Max processes/threads
10
11 # I/O limits
12 IOWeight=100           # I/O priority
```

Apply temporarily:

```
1 systemctl set-property nginx.service MemoryMax=1G
```

# Dependency Management

## Controlling Service Relationships

```
1 [Unit]
2 # Ordering: Start after these units
3 After=network.target postgresql.service
4
5 # Ordering: Start before these units
6 Before=nginx.service
7
8 # Requirements: Won't start without these (hard dependency)
9 Requires=postgresql.service
10
11 # Wants: Prefer these but can start without (soft dependency)
12 Wants=redis.service
13
14 # Conflicts: Can't run together
15 Conflicts=apache2.service
16
17 # Conditions: Only start if condition met
18 ConditionPathExists=/opt/myapp/config.yml
19 ConditionFileNotEmpty=/etc/myapp/config
```

# Debugging Failed Services

## Troubleshooting Checklist

```
1 # Check status and logs
2 systemctl status myapp.service
3 journalctl -u myapp.service -n 50
4
5 # Verify unit file
6 systemd-analyze verify /etc/systemd/system/myapp.service
7
8 # Test manually
9 sudo -u www-data /path/to/binary
10
11 # Check dependencies and permissions
12 systemctl list-dependencies myapp.service
13 ls -la /opt/myapp/
```

**Common Issues:** Wrong user/group, missing dependencies, incorrect paths, permissions

# Systemd Best Practices

## DO

- Use `Type=simple` for foreground apps
- Set appropriate `User=` and `Group=`
- Define clear dependencies
- Use `Restart=on-failure` for critical services
- Document with `Description=` and `Documentation=`
- Set resource limits
- Use `After=network.target` for network services
- Test unit files with `systemd-analyze verify`
- Use timers instead of cron
- Keep unit files in `/etc/systemd/system/`

## DON'T

- Run services as root unless necessary
- Use `Type=forking` for simple services
- Ignore failed dependencies
- Set `Restart=always` on oneshot services
- Edit files in `/lib/systemd/system/`
- Forget `daemon-reload` after changes
- Use absolute paths in `ExecStart` without testing
- Mix systemd and init.d scripts
- Ignore log messages
- Leave broken services enabled

# Real-World Examples

## Example 1: Node.js API Server

```
1 [Unit]
2 Description=Node.js API Server
3 After=network.target
4
5 [Service]
6 ExecStart=/usr/bin/node server.js
7 Restart=always
8
9 [Install]
10 WantedBy=multi-user.target
```

## Example 2: Database Backup Timer

```
1 # backup-db.service
2 [Service]
3 Type=oneshot
4 ExecStart=/usr/local/bin/backup.sh
5
6 # backup-db.timer
7 [Timer]
8 OnCalendar=*--*-* 03:00:00
```

# Quick Reference: systemctl

Command

Purpose

```
systemctl start/stop/restart SERVICE
```

Start, stop, or restart

```
systemctl status SERVICE
```

View status

```
systemctl enable/disable SERVICE
```

Start at boot or not

```
systemctl is-active SERVICE
```

Check if running

```
systemctl list-units --type=service
```

List services

```
systemctl --failed
```

List failed services

```
systemctl daemon-reload
```

Reload systemd config

```
systemctl cat SERVICE
```

Show unit file

# Quick Reference: journalctl

Command	Purpose
<code>journalctl</code> / <code>journalctl -f</code>	View all / follow logs
<code>journalctl -u SERVICE</code>	Logs for service
<code>journalctl -b</code> / <code>-k</code>	Boot logs / kernel
<code>journalctl -p err</code>	Only errors
<code>journalctl --since "1 hour ago"</code>	Time filter
<code>journalctl -n 50</code>	Last 50 lines
<code>journalctl --vacuum-time=7d</code>	Clean old logs

# Lab 3 Preview: Process & Service Management

## What You'll Do

1. Monitor processes with ps, top, and htop
2. Practice process control (signals, priorities)
3. Create a custom systemd service
4. Set up a systemd timer
5. Analyze boot performance
6. Troubleshoot failed services

## Deliverables

Process monitoring report, custom service file, timer config, boot analysis

**Time Estimate:** 2-3 hours

# Common Scenarios

## Scenario 1: High CPU Usage

```
1 # Find and investigate
2 top                                # Press 'P' to sort by CPU
3 ps aux --sort=-%cpu | head -10    # Top 10 consumers
4 lsof -p PID                         # What files?
5
6 # Fix
7 renice -n 10 -p PID
8 kill PID
```

## Scenario 2: Service Won't Start

```
1 systemctl status myapp.service
2 journalctl -u myapp.service -n 50
3 systemd-analyze verify /etc/systemd/system/myapp.service
```

## Scenario 3: Find Zombie Processes

```
1 ps aux | grep 'Z'
2 kill -9 PARENT_PID                  # Kill parent to reap
```

# Week 3 Action Items

## ✓ Before Next Lecture

1. Read **Chapter 10** (Processes)
2. Practice process monitoring
3. Experiment with systemctl on your VM
4. Create a simple custom service

## ✓ For Lab This Week

1. Complete **Lab 3: Process Management & Systemd**
2. Monitor processes, create services, set up timers
3. Debug service failures and optimize boot



## Practice

Find resource-hungry processes, create services, set up timers, debug configurations

# Key Takeaways

## Process Management

- Every program runs as a process with a unique PID
- Use `ps`, `top`, and `htop` to monitor processes
- Signals control process behavior - prefer SIGTERM over SIGKILL
- Nice values control CPU priority (-20 to 19)
- `/proc` filesystem exposes kernel information

## Systemd

- Modern init system and service manager
- Unit files define service behavior
- `systemctl` manages services, `journalctl` views logs
- Timers replace cron for scheduled tasks
- Resource limits prevent runaway services
- Dependencies ensure proper startup order

# Questions?



Processes and services are the heart of a running Linux system!

**Next Week:** Storage, Filesystems & LVM 

# Thank You!



**Remember:** Master the fundamentals - processes and services - to troubleshoot anything!

CSS 262 - Linux Administration & \*nix Systems for Cybersecurity