

# CSS 262: Linux Administration

## Storage, Filesystems & LVM

Lecture 4: Mastering Storage & Filesystems





# Today's Agenda

## Part 1: Storage Fundamentals

- Storage hierarchy overview
- Block devices & partitions
- Partition tables (MBR vs GPT)
- Disk management tools
- Device naming conventions

## Part 2: Filesystems & LVM

- Filesystem types & features
- Creating & mounting filesystems
- `/etc/fstab` configuration
- LVM architecture & benefits
- LVM operations & best practices

 **Learning Objective:** Understand Linux storage stack and implement flexible storage management with LVM.



# Quick Recap: Week 3

## Process Management & Systemd

- **Processes:** Fundamental execution units, managed via signals & priorities
- **Process States:** Running, sleeping, stopped, zombie
- **Key Tools:** `ps` , `top` , `htop` , `kill` , `nice` , `renice`
- **Systemd:** Modern init system and service manager
- **Unit Files:** Service definitions with dependencies & restart policies
- **journalctl:** Centralized logging system

 You should now be comfortable managing processes and creating systemd services!

# Part 1: Storage Fundamentals

Understanding the Linux Storage Stack



# The Storage Hierarchy

```
Error: Lexical error on line 3. Unrecognized text.  
...  B --> C1[/dev/sda1]      B --> C2[/dev/  
-----^
```

# Block Devices

Block devices are accessed in fixed-size blocks and allow random access.

## Device Naming Conventions

Type	Naming	Example
SATA/SCSI	/dev/sd[a-z]	/dev/sda
NVMe	/dev/nvme[0-9]n[0-9]	/dev/nvme0n1
Virtual	/dev/vd[a-z]	/dev/vda
LVM	/dev/mapper/	/dev/mapper/vg0-lv

## View Block Devices

```
1  lsblk      # Tree view  
2  lsblk -f    # Show filesystems
```



# Partition Tables: MBR vs GPT

## MBR (Master Boot Record)

**Legacy (1983)**

### Characteristics:

- Maximum 4 primary partitions
- Extended partitions for more logical partitions
- Max disk size: 2 TB
- 32-bit sector addressing
- Boot code in first sector

### Limitations:

- ❌ Small partition table
- ❌ No redundancy
- ❌ Limited disk size
- ❌ No partition names

## GPT (GUID Partition Table)

**Modern (2000s)**

### Characteristics:

- Up to 128 partitions (standard)
- Max disk size: 9.4 ZB
- 64-bit sector addressing
- CRC32 checksums for integrity
- Backup partition table at end

### Advantages:

- ✅ Large disk support
- ✅ Redundant partition table
- ✅ Partition names & GUIDs
- ✅ Future-proof design

# Partitioning Tools

## Command-Line Tools

### `fdisk` - Classic partition editor (MBR focus)

```
1  fdisk /dev/sda          # Interactive mode  
2  fdisk -l                # List all partitions  
3  fdisk -l /dev/sda        # List partitions on specific disk
```

### `gdisk` - GPT disk partitioner

```
1  gdisk /dev/sda          # Interactive GPT mode
```

### `parted` - Advanced partitioner (MBR & GPT)

```
1  parted /dev/sda print    # Show partition table  
2  parted /dev/sda mklabel gpt      # Create GPT table  
3  parted /dev/sda mkpart primary ext4 1MiB 100% # Create partition
```

**⚠ Warning:** Partitioning operations can destroy data! Always backup before making changes.



# Practical: Creating Partitions

Example: Create a new partition with `fdisk`

```
1 # Start fdisk
2 sudo fdisk /dev/sdb
3
4 # Commands within fdisk:
5 Command: n      # New partition
6 Partition: 1
7 First sector: [Enter]
8 Last sector: +10G
9
10 Command: t     # Change type
11 Hex code: 8e    # Linux LVM
12
13 Command: w     # Write changes
```

After Creating Partition

```
1 sudo partprobe /dev/sdb # Inform kernel
2 lsblk /dev/sdb          # Verify
```

# Part 2: Filesystems

Organizing Data on Storage



# What is a Filesystem?

A filesystem is a method of organizing and storing files on storage devices.

## Key Functions:

- **Data Organization:** Files, directories, metadata
- **Space Management:** Track free and used blocks
- **Access Control:** Permissions and ownership
- **Integrity:** Journaling, checksums, error detection
- **Performance:** Caching, buffering, optimization

## Filesystem Components:

- **Superblock:** Filesystem metadata (size, type, state)
- **Inodes:** File metadata (owner, permissions, timestamps)
- **Data Blocks:** Actual file content
- **Directory Entries:** Filename to inode mapping



# Linux Filesystem Types

Filesystem	Type	Key Features	Use Cases
<b>ext4</b>	Journaling	Mature, stable, widely supported	General purpose, root partition
<b>XFS</b>	Journaling	High performance, large files	Databases, media servers
<b>Btrfs</b>	Copy-on-Write	Snapshots, compression, RAID	Advanced features, snapshots
<b>F2FS</b>	Log-structured	Flash-optimized	SSDs, embedded systems
<b>NTFS</b>	Proprietary	Windows compatibility	Cross-platform storage
<b>FAT32/exFAT</b>	Simple	Universal compatibility	USB drives, SD cards
<b>ZFS</b>	Copy-on-Write	Enterprise features, integrity	High-end storage, NAS

 **Most Common:** ext4 (default), XFS (performance), Btrfs (features)

 **Tip:** Choose based on workload: ext4 for general, XFS for large files, Btrfs for snapshots



# Creating Filesystems

General Syntax: `mkfs.<type>`

```
1 # Create ext4 filesystem
2 sudo mkfs.ext4 /dev/sdb1
3
4 # Create ext4 with custom options
5 sudo mkfs.ext4 -L mydata -N 1000000 /dev/sdb1
6 # -L: Set volume label
7 # -N: Number of inodes
8
9 # Create XFS filesystem
10 sudo mkfs.xfs /dev/sdb1
11
12 # Create XFS with label
13 sudo mkfs.xfs -L backup /dev/sdb1
14
15 # Create Btrfs filesystem
16 sudo mkfs.btrfs /dev/sdb1
17
18 # Check filesystem after creation
19 sudo file -s /dev/sdb1
20 sudo blkid /dev/sdb1
```



**Caution:** `mkfs` destroys all data on the partition! Double-check the device name!



# Mounting Filesystems

**Mounting** attaches a filesystem to the directory tree.

## Manual Mounting

```
1 sudo mkdir -p /mnt/mydata
2 sudo mount /dev/sdb1 /mnt/mydata
3 df -h /mnt/mydata
4 sudo umount /mnt/mydata
```

## Mount Options

```
1 # Read-only
2 sudo mount -o ro /dev/sdb1 /mnt/mydata
3
4 # Security options
5 sudo mount -o noexec,nosuid /dev/sdb1 /mnt/mydata
```



# /etc/fstab

```
1 # <device>      <mount> <type> <options> <dump> <pass>
2 UUID=xxx-yyy   /       ext4    defaults  0        1
```

# 🔑 UUID vs Device Names

## Why Use UUIDs?

**Device names** (`/dev/sda1`) can change between reboots!

- Adding/removing drives
- Different boot order
- Driver load sequence

**UUIDs** (Universally Unique Identifiers) are persistent and unique.

## Finding UUIDs

```
1 # Method 1: blkid command
2 sudo blkid
3
4 # Method 2: lsblk with filesystem info
5 lsblk -f
6
7 # Method 3: By UUID directory
8 ls -l /dev/disk/by-uuid/
9
10 # Get specific device UUID
11 sudo blkid -s UUID -o value /dev/sdb1
```



# Practical: Setting Up /etc/fstab

## Step-by-Step Example

```
1 # Get UUID
2 UUID=$(sudo blkid -s UUID -o value /dev/sdb1)
3
4 # Create mount point
5 sudo mkdir -p /mnt/mydata
6
7 # Add to /etc/fstab
8 echo "UUID=$UUID /mnt/mydata ext4 defaults 0 2" | \
9   sudo tee -a /etc/fstab
10
11 # Test without rebooting
12 sudo mount -a
13
14 # Verify
15 df -h /mnt/mydata
```

**⚠️ Important:** Always test with `mount -a` before rebooting! Errors in fstab can prevent boot.



# Common Mount Options

## General Options

- `defaults` : Standard options (rw, suid, dev, exec, auto, nouser, async)
- `ro` / `rw` : Read-only / Read-write
- `auto` / `noauto` : Mount at boot / Don't mount automatically
- `user` / `nouser` : Allow user mounts / Root only
- `nofail` : Don't fail boot if device missing

## Security Options

- `noexec` : Prevent binary execution
- `nosuid` : Ignore setuid/setgid bits
- `nodev` : Don't interpret block/char devices

## Performance Options

- `noatime` : Don't update access time (faster)
- `nodiratime` : Don't update directory access time
- `relatime` : Update atime only if older than mtime
- `async` / `sync` : Async I/O / Sync I/O

## Example Combinations

```
1 # Web server data
2 defaults,noatime,noexec,nosuid
3
4 # User home directories
5 defaults,nodev,nosuid
6
7 # Backup drive (optional)
8 defaults,nofail,noatime
```

# Part 3: Logical Volume Management (LVM)

Flexible Storage Management



# What is LVM?

LVM (Logical Volume Manager) provides an abstraction layer between physical storage and filesystems.

## Traditional Partitioning Problems:

- X Fixed partition sizes
- X Difficult to resize
- X Can't span multiple disks easily
- X No snapshots

## LVM Benefits:

- ✓ Dynamic volume resizing (online!)
- ✓ Combine multiple disks into one volume
- ✓ Easy snapshots for backups
- ✓ Volume migration between disks
- ✓ Thin provisioning



**Real World:** LVM is standard in enterprise environments and modern Linux distributions.



# LVM Architecture

```
Error: Lexical error on line 2. Unrecognized text.  
...ph TD      A[/dev/sda] --> PV1[PV sda2]  
-----^
```



# LVM Components Explained

## Physical Volumes (PV)

- Physical disk/partition prepared for LVM
- Divided into **Physical Extents (PE)** (4 MB chunks)
- Commands: `pvcreate` , `pvs`

## Volume Groups (VG)

- Pool of physical volumes
- Commands: `vgcreate` , `vgs`

## Logical Volumes (LV)

- Virtual partitions from VG
- Can span PVs, resize, snapshot
- Commands: `lvcreate` , `lvs`



# Creating LVM: Step-by-Step

## 1. Create Physical Volumes

```
1 # Prepare partitions or disks for LVM
2 sudo pvcreate /dev/sdb1
3 sudo pvcreate /dev/sdc1
4
5 # Verify
6 sudo pvs
7 sudo pvdisplay
```

## 2. Create Volume Group

```
1 # Create VG from multiple PVs
2 sudo vgcreate vg_data /dev/sdb1 /dev/sdc1
3
4 # Verify
5 sudo vgs
6 sudo vgdisplay vg_data
```



# Creating LVM: Step-by-Step (cont.)

## 3. Create Logical Volumes

```
1 # Create LV with specific size
2 sudo lvcreate -L 50G -n lv_database vg_data
3
4 # Create LV with percentage of VG
5 sudo lvcreate -l 100%FREE -n lv_backups vg_data
6
7 # Verify
8 sudo lvs
9 sudo lvdisplay vg_data/lv_database
```

## 4. Create Filesystem and Mount

```
1 # Create filesystem on LV
2 sudo mkfs.ext4 /dev/vg_data/lv_database
3
4 # Mount it
5 sudo mkdir -p /mnt/database
6 sudo mount /dev/vg_data/lv_database /mnt/database
7
8 # Add to /etc/fstab for persistence
9 echo "/dev/vg_data/lv_database /mnt/database ext4 defaults 0 2" | sudo tee -a /etc/fstab
```



# Resizing Logical Volumes

## Extending a Logical Volume (Growing)

```
1 # Extend LV by 10 GB
2 sudo lvextend -L +10G /dev/vg_data/lv_database
3
4 # Resize filesystem (ext4)
5 sudo resize2fs /dev/vg_data/lv_database
6
7 # For XFS
8 sudo xfs_growfs /mnt/database
9
10 # Combined: extend + resize
11 sudo lvextend -r -L +10G /dev/vg_data/lv_database
```

### Requirements:

- VG must have free space
- Can be done while mounted!



# Shrinking Logical Volumes

## Reducing a Logical Volume (Shrinking)

```
1 sudo umount /mnt/database
2 sudo e2fsck -f /dev/vg_data/lv_database
3 sudo resize2fs /dev/vg_data/lv_database 40G
4 sudo lvreduce -L 40G /dev/vg_data/lv_database
5 sudo mount /dev/vg_data/lv_database /mnt/database
```

 **Warning:** Shrinking is dangerous! Backup first, unmount, check filesystem, shrink filesystem before LV. XFS does NOT support shrinking!



# LVM Snapshots

Snapshots create point-in-time copies of logical volumes.

## Creating & Using Snapshots

```
1 # Create snapshot
2 sudo lvcreate -L 5G -s -n lv_db_snap /dev/vg_data/lv_database
3
4 # Mount and backup
5 sudo mkdir -p /mnt/snapshot
6 sudo mount /dev/vg_data/lv_db_snap /mnt/snapshot
7 sudo tar czf backup.tar.gz -C /mnt/snapshot .
8
9 # Cleanup
10 sudo umount /mnt/snapshot
11 sudo lvremove /dev/vg_data/lv_db_snap
```



# LVM Information Commands

## Quick Status Commands

```
1 # Physical Volumes
2 pvs          # Summary
3 pvdisplay    # Detailed
4 pvdisplay /dev/sdb1 # Specific PV
5
6 # Volume Groups
7 vgs          # Summary
8 vgdisplay    # Detailed
9 vgdisplay vg_data # Specific VG
10
11 # Logical Volumes
12 lvs          # Summary
13 lvdisplay    # Detailed
14 lvdisplay vg_data/lv_database
```

## Advanced Information

```
1 # Show LV with more details
2 lvs -o +lv_size,lv_path,devices
3
4 # Show VG free space
5 vgs -o +vg_free,vg_size
6
7 # Show PV allocation
8 pvs -o +pv_used,pv_free
9
10 # Full system overview
11 sudo lsblk
12 sudo lsblk -f
```

 **Tip:** Add these to aliases: `alias lvss='sudo lvs -o +lv_size,devices'`



# Advanced LVM Operations

## Extending Volume Groups

```
1 sudo pvcreate /dev/sdd1
2 sudo vgextend vg_data /dev/sdd1
3 sudo vgs vg_data
```

## Moving Data Between PVs

```
1 sudo pvmove /dev/sdb1 /dev/sdd1
2 sudo vgreduce vg_data /dev/sdb1
3 sudo pvremove /dev/sdb1
```

## Renaming Volumes

```
1 sudo lvrename vg_data old_name new_name
2 sudo vgrename old_vg new_vg
```



# Storage Security Best Practices

## Mount Options for Security

```
1 # Uploads directory
2 /dev/vg_web/lv_uploads /var/www/uploads ext4 noexec,nosuid 0 2
```

## Encryption with LUKS

```
1 sudo cryptsetup luksFormat /dev/sdb1
2 sudo cryptsetup open /dev/sdb1 encrypted_disk
3 sudo pvcreate /dev/mapper/encrypted_disk
```

## Regular Backups & Monitoring

- Use LVM snapshots for consistent backups
- Monitor: `df -h`, `lvs`, `vgs`
- Check errors: `dmesg | grep -i error`



# Practical Lab Exercise

Scenario: Set up LVM for a web server

```
1 # Prepare disks
2 sudo pvcreate /dev/sdb1 /dev/sdc1
3 sudo vgcreate vg_webserver /dev/sdb1 /dev/sdc1
4
5 # Create logical volumes
6 sudo lvcreate -L 20G -n lv_www vg_webserver
7 sudo lvcreate -L 30G -n lv_mysql vg_webserver
8 sudo lvcreate -L 10G -n lv_logs vg_webserver
9
10 # Create filesystems
11 sudo mkfs.ext4 /dev/vg_webserver/lv_www
12 sudo mkfs.ext4 /dev/vg_webserver/lv_mysql
13 sudo mkfs.ext4 /dev/vg_webserver/lv_logs
14
15 # Mount
16 sudo mkdir -p /var/www /var/lib/mysql /var/log/apps
17 sudo mount /dev/vg_webserver/lv_www /var/www
18 sudo mount /dev/vg_webserver/lv_mysql /var/lib/mysql
19 sudo mount /dev/vg_webserver/lv_logs /var/log/apps
```

# Troubleshooting Storage Issues

## Common Issues

### Disk Full

```
1 sudo du -sh /* | sort -rh | head  
2 sudo journalctl --vacuum-size=100M
```

### LVM Issues

```
1 sudo vgchange -ay vg_data  
2 sudo pvscan  
3 sudo vgscan  
4 sudo lvscan
```

### Mount Failures

```
1 sudo mount -a  
2 sudo fsck /dev/sdb1
```

### Performance

```
1 iostat -x 1 5  
2 iotop
```



# Summary: Storage & Filesystems

## Key Concepts Covered

1. **Storage Hierarchy:** Physical disks → Partitions → Filesystems
2. **Partition Tables:** MBR (legacy) vs GPT (modern)
3. **Tools:** `fdisk` , `gdisk` , `parted` , `lsblk`
4. **Filesystems:** ext4, XFS, Btrfs - choose based on use case
5. **Mounting:** Manual ( `mount` ) and persistent( `/etc/fstab` )
6. **UUIDs:** Preferred over device names for stability

## LVM Benefits

- Flexible volume sizing (grow/shrink)
- Combine multiple disks
- Snapshots for backups
- Online operations
- Industry standard



# Learning Objectives: Did We Achieve?

By now, you should be able to:

- Understand Linux storage stack architecture
- Create and manage partitions with various tools
- Choose appropriate filesystem for different use cases
- Create filesystems and configure persistent mounts
- Explain LVM architecture (PV, VG, LV)
- Create and manage logical volumes
- Resize volumes and create snapshots
- Apply security best practices to storage
- Troubleshoot common storage issues

 **Next Week:** Bash Scripting & Automation - Automate all the things!

# Additional Resources

## Documentation

- LVM HOWTO - Comprehensive guide  
-----
- Red Hat Storage Guide  
-----
- Arch Wiki: LVM  
-----

## Books & Tools

- "*Linux Administration Handbook*" - Storage chapter
- `man lvm` , `man fstab` , `man mkfs`

## Practice

- Set up VMs with multiple disks
- Experiment with LVM resizing

# Questions?

Next: Bash Scripting & Automation

